

Comparing Nearest Neighbor, Clarke-Wright Savings, and Harmony Search for the Capacitated Vehicle Routing Problem

Mandy Jones, Tomu Yamashita, Tristen Martin, Charles Davis

Abstract

The Capacitated Vehicle Routing Problem (CVRP) is an optimization problem where the objective is to determine efficient routes that can satisfy customer demand without exceeding vehicle capacity constraints and ensure that each customer is visited exactly once. This study compares three algorithms to solving this problem: Nearest Neighbor, Clarke-Wright Savings, and Harmony Search. The performance of these algorithms is analyzed based on percentage gap from the best-known CVRPLIB solution, computational efficiency, and their ability to generate feasible routes for CVRP instances. Each algorithm is tested ten times on problem sizes ranging from 30 to 200 customers. The most notable finding was that Nearest Neighbor has the fastest runtimes with the least efficient routes; Clarke-Wright Savings had moderate speed with good routes; and Harmony Search generated the most efficient routes with the slowest run time.

1. Problem Definition

The Vehicle Routing Problem (VRP) seeks the optimal set of routes for a fleet of vehicles delivering goods to customers while minimizing total travel distance. This study focuses on the Capacitated Vehicle Routing Problem (CVRP), where each vehicle has a fixed carrying capacity that limits how much demand it can serve per route, making the problem harder to solve. The CVRP is NP-hard, since it reduces to the Traveling Salesman Problem when using a single, unlimited-capacity vehicle, meaning no known algorithm can guarantee an optimal solution in polynomial time, which is why heuristic approaches are used instead.

2. Algorithm Explanations

2.1 Nearest Neighbor Heuristic

In this comparison study, the Nearest Neighbor algorithm serves as the baseline. This is because it is the simplest approach: it performs the closest thing it hasn't yet. To explain, it starts at the depot, which is the starting and ending point for the truck's route. It then goes to the unvisited customer that is the closest and fits into the truck's remaining space. The algorithm repeats this until no more customers fit, and then it sends the truck home and starts a new one. It does this until every customer is visited.

2.2 Clarke-Wright Savings Algorithm

The Clarke-Wright Savings algorithm starts by giving every customer their own truck, having the truck go from the depot to that customer and back. It then looks at every possible pair of customers and calculates the savings, which is how much distance would be saved by combining their two separate trips into one that visits both. These savings values are sorted from largest to smallest, with the algorithm merging routes together starting with the largest savings. A merge only happens if the two customers are the first or last customer in the list of their routes and if combining them doesn't exceed the truck's capacity. This process continues until no more valid merges are left.

2.3 Harmony Search

The Harmony Search algorithm starts by creating a harmony memory, which stores several possible customer orderings along with the total distance of the truck routes they produce. It begins with multiple different solutions so that the search can explore several possibilities instead of starting from a single route. Unlike the traditional Harmony Search algorithm, this implementation uses deterministic rules rather than random selection to perform memory consideration and pitch adjustment, ensuring the same input always produces the same output. It then creates new customer orderings by building a sequence of one customer at a time, either using information from existing solutions in memory (memory consideration) or selecting a different available customer. The algorithm can also adjust choices by moving toward nearby customers (pitch adjustment) to improve the route. Each new ordering is converted into feasible truck routes and refined by swapping neighboring customers when those swaps reduce the total distance. If the improved solution is better than the worst solution currently stored in memory, it replaces that solution. This process repeats a set number of iterations, and the best solution remaining in the harmony memory is returned.

3. Theoretical Complexity

3.1 Nearest Neighbor Heuristic

The Nearest Neighbor algorithm has a time complexity of $\Theta(n^2)$. Every time a customer is added to a route, the algorithm must scan all unvisited customers to identify the next closest one.

3.2 Clarke-Wright Savings Algorithm

Clarke-Wright Savings algorithm has a time complexity of $\Theta(n^2 \log n)$. Calculating the savings value for every pair of customers requires comparing all possible pairs, which takes $\Theta(n^2)$ time. The n^2 savings values must then be sorted from largest to smallest so that the most beneficial merges are attempted first. This sorting step takes $\Theta(n^2 \log n)$ time and is the most expensive.

3.3 Harmony Search

Harmony Search has a time complexity of $\Theta(I n^3)$, where I is the number of iterations and n is the number of customers. During each iteration, the algorithm generates and evaluates a new customer ordering before improving it using local search. Constructing each ordering requires repeatedly searching through and removing customers from lists, while the local search evaluates many neighboring orderings to find improvements. Together, these operations require $\Theta(n^3)$ time per iteration. Repeating this process for I iterations gives an overall time complexity of $\Theta(I n^3)$.

4. Experimental Design

All three algorithms were implemented in Python. Inputs were four real benchmark instances from CVRPLIB (A-n32-k5, E-n51-k5, E-n101-k8, M-n200-k16), giving problem sizes of 31, 50, 100, and 199 customers with fixed vehicle capacities defined by each instance. These are established, non-random datasets with published best-known solutions, which serve as our reference method, since brute force is infeasible at these sizes. Each algorithm was run ten times per instance, recording route distance and runtime for each trial. Performance metrics were average execution time in milliseconds and percentage gap from the best-known CVRPLIB solution. All results were written to "results/all_results.csv". Experiments were run on various computers including Windows 11.

Our research questions focus on the trade-off between solution quality and computational cost. Because the CVRP is NP-hard, the practical question becomes which heuristic gives the best value for its runtime. We expected Nearest Neighbor to be fastest but least accurate due to its greedy, irreversible decisions; Clarke-Wright to offer a middle ground by evaluating all pairwise savings upfront; and Harmony Search to produce the best routes at the highest computational cost, since metaheuristics explore many candidate solutions. Question 3 was chosen to test whether any algorithm degrades at scale, which matters for real-world applicability.

The experiment was structured around three research questions:

1. How does runtime change as the number of customers increases?
2. Which algorithm provides the best balance between percentage gap from the best-known CVRPLIB solution and execution time?
3. Does each algorithm's percentage gap from the best-known solution stay consistent as problem size increases, or does it get better or worse?

5. Results & Analysis

Nearest Neighbor Results:

data file	num_customers	num_routes	total_distance	best value	gap_percent	avg_runtime_ms	trials
data/A-n32-k5.vrp	31	5	1146.4	784	46.22	0.5312	10
data/E-n51-k5.vrp	50	5	711.5	521	36.56	1.1336	10
data/E-n101-k8.vrp	100	8	1174.36	815	44.09	3.3567	10
data/M-n200-k16.vrp	199	17	2034.4	N/A	N/A	16.4016	10

As the problem grew from 31 customers to 199 customers, the runtime jumped from 0.53ms to 16.40ms, which is 31 times slower. That makes sense regarding the $\Theta(n^2)$ complexity, since the algorithm must compare every remaining customer at each step. The bigger the problem, the more comparisons that have to be made. However, the quality of the routes stayed in the same range (36% - 46%) as the problem grew, which makes sense because the algorithm doesn't reconsider earlier choices. Its greedy decision-making causes the same degree of inefficiency regardless of problem size.

Clarke-Wright Savings Results:

clarke_wright	data/A-n32-k5.vrp	31	5	843.69	784	7.61	1.1698	10
clarke_wright	data/E-n51-k5.vrp	50	6	584.64	521	12.21	3.1485	10
clarke_wright	data/E-n101-k8.vrp	100	8	889.0	815	9.08	16.0064	10
clarke_wright	data/M-n200-k16.vrp	199	17	1395.74	N/A	N/A	130.4101	10

As the problem grew from 31 to 199 customers, the runtime jumped from 1.17ms to 130.41ms, which is 111 times slower. This is much steeper than shown in the Nearest Neighbor's growth on the same scale. This makes sense as Clarke-Wright Savings' time complexity is $\Theta(n^2 \log n)$; it grows a bit faster because of the extra step of sorting the saving values. Meanwhile, the gap from the best value stayed in a lower, less expansive range from 7% - 12%, since the algorithm considers all possible customer pairings upfront rather than greedily picking what's closest.

Harmony Search Results:

harmony_search	data/A-n32-k5.vrp	31	5	1714.2	784	118.65	155.3155	10
harmony_search	data/E-n51-k5.vrp	50	6	1281.13	521	145.9	416.1114	10
harmony_search	data/E-n101-k8.vrp	100	8	2106.33	815	158.45	2204.2205	10
harmony_search	data/M-n200-k16.vrp	199	17	4915.27	N/A	N/A	49263.7119	10

As the problem grew from 31 to 199 customers, the runtime jumped from 155.32ms to 49263.71ms, which is 317 times slower. This is the steepest growth yet, compared to both Nearest Neighbor as well as Clarke-Wright. Since Harmony Search's time complexity is $\Theta(n^3)$, it makes sense that the average runtime is much larger. The gap from the best value was

consistently more than double, ranging between 119% - 158%. This large gap suggests that the additional time was spent searching for a better path.

Looking at the results, the empirical analysis directly answers the three research questions. For Question 1, runtime increased with problem size for all three algorithms, and the relative steepness of that growth matched the expected ordering of their complexity classes: Nearest Neighbor grew the slowest, Clarke-Wright faster, and Harmony Search the fastest by a wide margin. For Question 2, Clarke-Wright offered the best balance. It produced the lowest gaps of the three while running faster than Harmony Search, and its extra runtime relative to Nearest Neighbor bought a substantial gain in solution quality. For Question 3, the answer differs by algorithm: Nearest Neighbor and Clarke-Wright held roughly stable gaps across the instances with best-known solutions, while Harmony Search degraded steadily as problem size grew. With only three instances having best-known solutions, trend claims about gap percentage are suggestive rather than conclusive.

In terms of limitations, Harmony Search's large gap comes down to our deterministic implementation. Without randomized selection, the ten starting harmonies are just rotations and reversals of the same sorted customer list, and every new harmony is a fixed function of the iteration number, so the search never explores much of the solution space. Our parameter choices didn't help either: the iteration count was fixed at 200 regardless of instance size, which likely explains why the gap grew steadily as the problems got larger. A well-tuned stochastic version with a search budget that scales with problem size would almost certainly do better than what we measured. We also only tested four instances, and only three of those had published best-known values, so we can't say much with confidence about the trends. And since we measured wall-clock time, our numbers are tied to the machine we ran them on; counting basic operations would have been machine-independent. If we did this again, we'd tune the Harmony Search parameters systematically, test multiple instances at each size, and track operation counts alongside runtime. That leaves an open question worth following up on: how much of Harmony Search's gap could tuning and randomized selection close, and would the extra runtime ever be worth it compared to Clarke-Wright?

6. References

[K-Nearest Neighbor\(KNN\) Algorithm - GeeksforGeeks](#)

[Solving CVRP with Greedy Algorithms](#)

<https://github.com/ngchunlong279645/Capacitated-Vehicle-Routing-Problem-CVRP->

<https://github.com/zuzhaoye/vehicle-routing-problem-vrp-Clarke-Wright-Savings-Method>

[Improved Harmony Search Algorithm](#)

[Critical Analysis on Harmony Search Algorithm](#)

https://github.com/sagniklp/Harmony-Search?utm_source=chatgpt.com